

# Group Exercise 2 (Ungraded)

## Data Importing

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import StandardScaler
```

## Importing datasets

```
In [3]: retail_orders = pd.read_csv("Retail_Data_Orders_W23.csv") # Imported Retail_Data_Or
retail_data = pd.read_csv("Retail_Data_W23.csv") # Imported Retail_Data_W23.csv
retail_unlabelled_data = pd.read_csv("Retail_Unlabeled_Data.csv") # Imported Retail
store_data = pd.read_csv("Store.csv") # Imported Store.csv

retail_orders_df = pd.DataFrame(retail_orders) # Created a dataframe for retail_or
retail_data_df = pd.DataFrame(retail_data) # Created a dataframe for retail_data.cs
retail_unlabelled_data_df = pd.DataFrame(retail_unlabelled_data) # Created a datafr
store_data_df = pd.DataFrame(store_data) # Created a dataframe for store_data.csv
```

## Exploratory Data Analysis

```
In [4]: retail_orders_df.head() # Top 5 rows of the dataframe
retail_orders_df.shape # Shape of the dataframe
```

```
Out[4]: (651013, 3)
```

```
In [6]: retail_data_df.head() # Top 5 rows of the dataframe
retail_data_df.shape # Shape of the dataframe
```

```
Out[6]: (651013, 10)
```

```
In [7]: retail_unlabelled_data_df.head() # Top 5 rows of the dataframe
retail_unlabelled_data_df.shape # Shape of the dataframe
```

```
Out[7]: (162754, 10)
```

```
In [8]: store_data_df.head() # Top 5 rows of the dataframe
store_data_df.shape # Shape of the dataframe
```

```
Out[8]: (1115, 10)
```

## Cleaning and Pre-processing

Our first step is to merge the three datasets into one and clean them.

Data sets to merge:

1. retail\_orders\_df
2. retail\_data\_df
3. store\_data\_df

```
In [9]: merged_1 = retail_data_df.merge(retail_orders_df, on="Id", how="inner")
```

```
In [10]: merged_1.head()
```

```
Out[10]:
```

	Unnamed: 0_x	Store	DayOfWeek	Date	Customers	Open	Promo	StateHoliday	SchoolHoliday
0	339662	516	5	2014-09-12	665	1	0	0	0
1	969929	665	2	2013-02-12	1213	1	0	0	0
2	499245	511	4	2014-04-10	685	1	0	0	1
3	581706	462	7	2014-01-26	0	0	0	0	0
4	618037	1113	3	2013-12-25	0	0	0	c	1

```
In [11]: merged_final = merged_1.merge(store_data_df, on="Store", how="left")
```

```
In [12]: merged_final.head()
```

```
Out[12]:
```

	Unnamed: 0_x	Store	DayOfWeek	Date	Customers	Open	Promo	StateHoliday	SchoolHoliday
0	339662	516	5	2014-09-12	665	1	0	0	0
1	969929	665	2	2013-02-12	1213	1	0	0	0
2	499245	511	4	2014-04-10	685	1	0	0	1
3	581706	462	7	2014-01-26	0	0	0	0	0
4	618037	1113	3	2013-12-25	0	0	0	c	1

5 rows × 21 columns

```
In [14]: merged_final.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 651013 entries, 0 to 651012
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Unnamed: 0_x                          651013 non-null int64
1   Store                                651013 non-null int64
2   DayOfWeek                            651013 non-null int64
3   Date                                 651013 non-null object
4   Customers                            651013 non-null int64
5   Open                                 651013 non-null int64
6   Promo                                651013 non-null int64
7   StateHoliday                         651013 non-null object
8   SchoolHoliday                        651013 non-null int64
9   Id                                    651013 non-null int64
10  Unnamed: 0_y                          651013 non-null int64
11  Orders                               651013 non-null int64
12  StoreType                            651013 non-null object
13  Assortment                           651013 non-null object
14  CompetitionDistance                  649312 non-null float64
15  CompetitionOpenSinceMonth            443851 non-null float64
16  CompetitionOpenSinceYear             443851 non-null float64
17  Promo2                               651013 non-null int64
18  Promo2SinceWeek                      325567 non-null float64
19  Promo2SinceYear                      325567 non-null float64
20  PromoInterval                       325567 non-null object
dtypes: float64(5), int64(11), object(5)
memory usage: 109.3+ MB

```

```
In [15]: merged_final = merged_final.drop(['Unnamed: 0_y', 'Unnamed: 0_x'], axis=1)
```

```
In [17]: merged_final.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 651013 entries, 0 to 651012
Data columns (total 19 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Store                                651013 non-null int64
1   DayOfWeek                            651013 non-null int64
2   Date                                 651013 non-null object
3   Customers                            651013 non-null int64
4   Open                                 651013 non-null int64
5   Promo                                651013 non-null int64
6   StateHoliday                         651013 non-null object
7   SchoolHoliday                        651013 non-null int64
8   Id                                    651013 non-null int64
9   Orders                               651013 non-null int64
10  StoreType                            651013 non-null object
11  Assortment                           651013 non-null object
12  CompetitionDistance                  649312 non-null float64
13  CompetitionOpenSinceMonth            443851 non-null float64
14  CompetitionOpenSinceYear             443851 non-null float64
15  Promo2                               651013 non-null int64
16  Promo2SinceWeek                      325567 non-null float64
17  Promo2SinceYear                      325567 non-null float64
18  PromoInterval                       325567 non-null object
dtypes: float64(5), int64(9), object(5)
memory usage: 99.3+ MB

```

```
In [18]: merged_final.isnull().sum()
```

```
Out[18]: Store 0
DayOfWeek 0
Date 0
Customers 0
Open 0
Promo 0
StateHoliday 0
SchoolHoliday 0
Id 0
Orders 0
StoreType 0
Assortment 0
CompetitionDistance 1701
CompetitionOpenSinceMonth 207162
CompetitionOpenSinceYear 207162
Promo2 0
Promo2SinceWeek 325446
Promo2SinceYear 325446
PromoInterval 325446
dtype: int64
```

```
In [19]: merged_final.CompetitionDistance.describe()
```

```
Out[19]: count    649312.000000
mean      5436.342390
std       7713.881629
min       20.000000
25%      710.000000
50%     2330.000000
75%     6890.000000
max     75860.000000
Name: CompetitionDistance, dtype: float64
```

```
In [20]: merged_final['CompetitionDistance'].fillna(merged_final['CompetitionDistance'].medi
```

```
In [22]: merged_final['CompetitionOpenSinceMonth'].fillna(0, inplace=True)
merged_final['CompetitionOpenSinceYear'].fillna(0, inplace=True)
```

```
In [24]: merged_final['Promo2SinceWeek'].fillna(0, inplace=True)
merged_final['Promo2SinceYear'].fillna(0, inplace=True)
```

```
In [25]: merged_final['Promo2SinceWeek'].fillna(0, inplace=True)
merged_final['Promo2SinceYear'].fillna(0, inplace=True)
```

```
In [26]: merged_final['PromoInterval'].fillna('PromoUnavailable', inplace=True)
```

```
In [27]: merged_final.isnull().sum()
```

```
Out[27]: Store 0
DayOfWeek 0
Date 0
Customers 0
Open 0
Promo 0
StateHoliday 0
SchoolHoliday 0
Id 0
Orders 0
StoreType 0
Assortment 0
CompetitionDistance 0
CompetitionOpenSinceMonth 0
CompetitionOpenSinceYear 0
Promo2 0
Promo2SinceWeek 0
Promo2SinceYear 0
PromoInterval 0
dtype: int64
```

```
In [28]: merged_final.describe()
```

```
Out[28]:
```

	Store	DayOfWeek	Customers	Open	Promo	SchoolHoliday
<b>count</b>	651013.000000	651013.000000	651013.000000	651013.000000	651013.000000	651013.000000
<b>mean</b>	558.645629	3.999336	632.851832	0.829619	0.381558	0.178927
<b>std</b>	321.905872	1.998260	464.857658	0.375967	0.485769	0.383292
<b>min</b>	1.000000	1.000000	0.000000	0.000000	0.000000	0.000000
<b>25%</b>	281.000000	2.000000	404.000000	1.000000	0.000000	0.000000
<b>50%</b>	558.000000	4.000000	609.000000	1.000000	0.000000	0.000000
<b>75%</b>	838.000000	6.000000	837.000000	1.000000	1.000000	0.000000
<b>max</b>	1115.000000	7.000000	5458.000000	1.000000	1.000000	1.000000

```
In [29]: merged_final.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 651013 entries, 0 to 651012
Data columns (total 19 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Store                                651013 non-null  int64
1   DayOfWeek                            651013 non-null  int64
2   Date                                651013 non-null  object
3   Customers                            651013 non-null  int64
4   Open                                651013 non-null  int64
5   Promo                                651013 non-null  int64
6   StateHoliday                         651013 non-null  object
7   SchoolHoliday                       651013 non-null  int64
8   Id                                    651013 non-null  int64
9   Orders                              651013 non-null  int64
10  StoreType                            651013 non-null  object
11  Assortment                           651013 non-null  object
12  CompetitionDistance                  651013 non-null  float64
13  CompetitionOpenSinceMonth            651013 non-null  float64
14  CompetitionOpenSinceYear             651013 non-null  float64
15  Promo2                               651013 non-null  int64
16  Promo2SinceWeek                      651013 non-null  float64
17  Promo2SinceYear                      651013 non-null  float64
18  PromoInterval                       651013 non-null  object
dtypes: float64(5), int64(9), object(5)
memory usage: 99.3+ MB

```

```

In [32]: merged_final.CompetitionDistance = merged_final.CompetitionDistance.astype('int64')
merged_final.CompetitionOpenSinceMonth = merged_final.CompetitionOpenSinceMonth.as
merged_final.CompetitionOpenSinceYear = merged_final.CompetitionOpenSinceYear.as
merged_final.Promo2SinceWeek = merged_final.Promo2SinceWeek.astype('int').astype('c
merged_final.Promo2SinceYear = merged_final.Promo2SinceYear.astype('int').astype('c
merged_final.DayOfWeek = merged_final.DayOfWeek.astype('int').astype('object')

```

```

In [33]: merged_final.info()

```

```


<class 'pandas.core.frame.DataFrame'>
Int64Index: 651013 entries, 0 to 651012
Data columns (total 19 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Store                                651013 non-null  int64
1   DayOfWeek                            651013 non-null  object
2   Date                                651013 non-null  object
3   Customers                            651013 non-null  int64
4   Open                                651013 non-null  int64
5   Promo                                651013 non-null  int64
6   StateHoliday                         651013 non-null  object
7   SchoolHoliday                       651013 non-null  int64
8   Id                                    651013 non-null  int64
9   Orders                              651013 non-null  int64
10  StoreType                            651013 non-null  object
11  Assortment                           651013 non-null  object
12  CompetitionDistance                  651013 non-null  int64
13  CompetitionOpenSinceMonth            651013 non-null  object
14  CompetitionOpenSinceYear             651013 non-null  object
15  Promo2                               651013 non-null  int64
16  Promo2SinceWeek                      651013 non-null  object
17  Promo2SinceYear                      651013 non-null  object
18  PromoInterval                       651013 non-null  object
dtypes: int64(9), object(10)
memory usage: 99.3+ MB

```

```
In [34]: merged_final.drop(['Date'],axis = 1,inplace = True)
```

```
In [ ]:
```


```
In [35]: plt.hist(merged_final['Orders'])  
plt.xlabel('Orders')  
plt.ylabel('Frequency')  
plt.title('Distribution of Orders')  
plt.show()
```



```
In [ ]:
```


```
In [36]: plt.figure(figsize=(8,5))  
sns.heatmap(merged_final.corr(), annot=True, cmap='coolwarm', linewidths=0.5)  
plt.title('Correlation plot')  
plt.show()
```

C:\Users\balde\AppData\Local\Temp\ipykernel\_15332\2851571332.py:2: FutureWarning:  
The default value of numeric\_only in DataFrame.corr is deprecated. In a future version,  
it will default to False. Select only valid columns or specify the value of  
numeric\_only to silence this warning.  
sns.heatmap(merged\_final.corr(), annot=True, cmap='coolwarm', linewidths=0.5)



### Scatter plot comparison between Orders and Customers


```
In [55]: sns.regplot(data=merged_final, x = 'Customers', y = 'Orders', line_kws=dict(color='red'))
plt.title('Scatter plot with regression line')
plt.show()
```





## Scatter plot comparison between Customers and Orders

```
In [56]: sns.regplot(data=merged_final, x = 'Orders', y = 'Customers', line_kws=dict(color="r"))  
plt.title('Scatter plot with regression line')  
plt.show()
```




In [ ]:

In [ ]:

## Scatter plot comparison between Store and Orders

```
In [42]: sns.regplot(data=merged_final, x = 'Orders', y = 'Store')  
plt.title('Scatter plot with regression line')  
plt.show()
```



In [ ]:

```
In [43]: merged_final.select_dtypes(include='object').columns
```

```
Out[43]: Index(['DayOfWeek', 'StateHoliday', 'StoreType', 'Assortment',  
          'CompetitionOpenSinceMonth', 'CompetitionOpenSinceYear',  
          'Promo2SinceWeek', 'Promo2SinceYear', 'PromoInterval'],  
          dtype='object')
```

```
In [44]: df_encoded = pd.get_dummies(merged_final, columns = ['StateHoliday', 'StoreType', 'Assortment',  
          'Promo2SinceWeek', 'Promo2SinceYear', 'PromoInterval'])  
df_encoded.columns
```

```

C:\Users\balde\AppData\Local\Temp\ipykernel_15332\475200570.py:1: FutureWarning: In
a future version, the Index constructor will not infer numeric dtypes when passe
d object-dtype sequences (matching Series behavior)
    df_encoded = pd.get_dummies(merged_final, columns = ['StateHoliday','StoreTyp
e','Assortment','CompetitionOpenSinceMonth','CompetitionOpenSinceYear',
C:\Users\balde\AppData\Local\Temp\ipykernel_15332\475200570.py:1: FutureWarning: In
a future version, the Index constructor will not infer numeric dtypes when passe
d object-dtype sequences (matching Series behavior)
    df_encoded = pd.get_dummies(merged_final, columns = ['StateHoliday','StoreTyp
e','Assortment','CompetitionOpenSinceMonth','CompetitionOpenSinceYear',
C:\Users\balde\AppData\Local\Temp\ipykernel_15332\475200570.py:1: FutureWarning: In
a future version, the Index constructor will not infer numeric dtypes when passe
d object-dtype sequences (matching Series behavior)
    df_encoded = pd.get_dummies(merged_final, columns = ['StateHoliday','StoreTyp
e','Assortment','CompetitionOpenSinceMonth','CompetitionOpenSinceYear',
C:\Users\balde\AppData\Local\Temp\ipykernel_15332\475200570.py:1: FutureWarning: In
a future version, the Index constructor will not infer numeric dtypes when passe
d object-dtype sequences (matching Series behavior)
    df_encoded = pd.get_dummies(merged_final, columns = ['StateHoliday','StoreTyp
e','Assortment','CompetitionOpenSinceMonth','CompetitionOpenSinceYear',

```

```

Out[44]: Index(['Store', 'Customers', 'Open', 'Promo', 'SchoolHoliday', 'Id', 'Orders',
      'CompetitionDistance', 'Promo2', 'StateHoliday_0',
      ...
      'PromoInterval_Jan, Apr, Jul, Oct', 'PromoInterval_Mar, Jun, Sept, Dec',
      'PromoInterval_PromoUnavailable', 'DayOfWeek_1', 'DayOfWeek_2',
      'DayOfWeek_3', 'DayOfWeek_4', 'DayOfWeek_5', 'DayOfWeek_6',
      'DayOfWeek_7'],
      dtype='object', length=101)

```

```

In [45]: # Model fitting and splitting of dataset

target_column = 'Orders'
X = df_encoded.drop(['Orders'],axis = 1)
y = df_encoded['Orders']

# Train Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.75, test_size=0.25)

# Scaling is done after train test split to prevent data leakage
scaler = StandardScaler()
X_train_sc = scaler.fit_transform(X_train)
X_test_sc = scaler.transform(X_test)

model = LinearRegression()
model.fit(X_train_sc, y_train)

# Make predictions on the scaled test data
ypred = model.predict(X_test_sc)

```

```

In [46]: # Evaluating Mean squared error and R2 scores

mse = mean_squared_error(y_test, ypred)
r2 = r2_score(y_test, ypred)

# Print evaluation metrics

```

```
print(f"Mean Squared Error : {mse}")
print(f"R-squared : {r2}")
```

```
Mean Squared Error : 1356489.415316727
R-squared : 0.9082723261670267
```

In [ ]:

In [ ]:

## Build an OLS model in Excel/Python on select features of the Retail dataset without using in-built OLS functions

```
In [47]: import pandas as pd
import numpy as np
```

```
In [48]: def ordinary_least_squares(X, Y):
    # Step 1: Calculate the means of X and Y
    mean_X = np.mean(X)
    mean_Y = np.mean(Y)

    # Step 2: Calculate the deviations from the means
    dev_X = X - mean_X
    dev_Y = Y - mean_Y

    # Step 3: Calculate the slope (m)
    m = np.sum(dev_X * dev_Y) / np.sum(dev_X ** 2)

    # Step 4: Calculate the intercept (b)
    b = mean_Y - m * mean_X

    # Step 5: Return the coefficients
    return m, b
```

```
In [49]: def calculate_r_squared(X, Y, slope, intercept):
    # Step 1: Calculate the predicted values (Y_pred) using the linear regression model
    Y_pred = slope * X + intercept

    # Step 2: Calculate the total sum of squares (TSS)
    mean_Y = np.mean(Y)
    tss = np.sum((Y - mean_Y) ** 2)

    # Step 3: Calculate the residual sum of squares (RSS)
    rss = np.sum((Y - Y_pred) ** 2)

    # Step 4: Calculate R-squared (R²)
    r_squared = 1 - (rss / tss)

    return r_squared
```

```
In [53]: x = merged_final['Orders'].values
y = merged_final['Customers'].values

# Step 6: Use the function to get coefficients for your data

slope, intercept = ordinary_least_squares(x, y)
r_squared = calculate_r_squared(x, y, slope, intercept)

# Step 7: Output the results
print("Slope (m):", slope)
```

```
print("Intercept (b):", intercept)
print("R-squared (R²):", r_squared)
```

Slope (m): 0.107955007903139  
Intercept (b): 9.8951915907096  
R-squared (R²): 0.7998921458551951

```
In [57]: import seaborn as sns
import matplotlib.pyplot as plt
sns.regplot(x=x, y=y, ci=None, line_kws=dict(color="r"))

# Show the plot
plt.xlabel('X')
plt.ylabel('Y')
plt.title('Scatterplot with Regression Line (regplot)')
plt.show()
```



In [ ]: